

Times Series Econometrics Project

```
In [1]: # Import Libraries
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
sns.set(style="whitegrid")

from pandas_datareader import data as web
import datetime

from arch.unitroot import ADF, KPSS
from statsmodels.tsa.vector_ar.vecm import coint_johansen, VECM
from statsmodels.tsa.api import VAR

import warnings
warnings.filterwarnings("ignore")
```

Data gathering and collection

```
In [2]: # Data gathering and collection

start = datetime.datetime(1959, 1, 1)
end = datetime.datetime(2025, 10, 1)

# Download data from FRED
pce = web.DataReader('PCECCC96', 'fred', start, end)
dpi = web.DataReader('DPIC96', 'fred', start, end)
gdp = web.DataReader('GDP', 'fred', start, end) # optional variable for advanced work

# Merge and clean
data = pd.concat([pce, dpi, gdp], axis=1)
data.columns = ['PCE', 'DPI', 'GDP']
data = data.dropna()
data.tail()
```

Out[2]:

	PCE	DPI	GDP
DATE			
2024-04-01	16009.637	17700.963	29147.044
2024-07-01	16165.768	17755.291	29511.664
2024-10-01	16320.890	17843.165	29825.182
2025-01-01	16345.793	17943.157	30042.113
2025-04-01	16445.685	18082.022	30485.729

Pre-Estimation Analysis

In [3]:

```
# Pre-Estimation Analysis

# Visual inspection
data.plot(figsize=(12,6), title="PCE, DPI, GDP over time")
plt.show()


# Log transformation for variance stabilization
data['ln_pce'] = np.log(data['PCE'])
data['ln_dpi'] = np.log(data['DPI'])
data['ln_gdp'] = np.log(data['GDP'])

# Stationarity tests
def adf_test(series, title):
    print(f'ADF Test: {title}')
    print(ADF(series, trend='ct').summary())
    print("\n")

def kpss_test(series, title):
    print(f'KPSS Test: {title}')
    print(KPSS(series, trend='ct').summary())
    print("\n")

print("== Stationarity Tests ==")
adf_test(data['ln_pce'], 'ln_pce level')
```

```
adf_test(data['ln_dpi'], 'ln_dpi level')
adf_test(data['ln_gdp'], 'ln_gdp level')
adf_test(data['ln_pce'].diff().dropna(), 'Δln_pce')
adf_test(data['ln_dpi'].diff().dropna(), 'Δln_dpi')
adf_test(data['ln_gdp'].diff().dropna(), 'Δln_gdp')
```



```
==== Stationarity Tests ====
ADF Test: ln_pce level
    Augmented Dickey-Fuller Results
=====
Test Statistic          -1.564
P-value                0.806
Lags                  0
-----
```

Trend: Constant and Linear Time Trend
Critical Values: -3.99 (1%), -3.43 (5%), -3.14 (10%)
Null Hypothesis: The process contains a unit root.
Alternative Hypothesis: The process is weakly stationary.

```
ADF Test: ln_dpi level
    Augmented Dickey-Fuller Results
=====
Test Statistic          -2.302
P-value                0.433
Lags                  6
-----
```

Trend: Constant and Linear Time Trend
Critical Values: -3.99 (1%), -3.43 (5%), -3.14 (10%)
Null Hypothesis: The process contains a unit root.
Alternative Hypothesis: The process is weakly stationary.

```
ADF Test: ln_gdp level
    Augmented Dickey-Fuller Results
=====
Test Statistic          -0.399
P-value                0.987
Lags                  2
-----
```

Trend: Constant and Linear Time Trend
Critical Values: -3.99 (1%), -3.43 (5%), -3.14 (10%)
Null Hypothesis: The process contains a unit root.
Alternative Hypothesis: The process is weakly stationary.

```
ADF Test: Δln_pce
    Augmented Dickey-Fuller Results
=====
Test Statistic      -17.101
P-value            0.000
Lags               0
-----
```

Trend: Constant and Linear Time Trend
Critical Values: -3.99 (1%), -3.43 (5%), -3.14 (10%)
Null Hypothesis: The process contains a unit root.
Alternative Hypothesis: The process is weakly stationary.

```
ADF Test: Δln_dpi
    Augmented Dickey-Fuller Results
=====
Test Statistic      -8.423
P-value            0.000
Lags               5
-----
```

Trend: Constant and Linear Time Trend
Critical Values: -3.99 (1%), -3.43 (5%), -3.14 (10%)
Null Hypothesis: The process contains a unit root.
Alternative Hypothesis: The process is weakly stationary.

```
ADF Test: Δln_gdp
    Augmented Dickey-Fuller Results
=====
Test Statistic      -9.169
P-value            0.000
Lags               1
-----
```

Trend: Constant and Linear Time Trend
Critical Values: -3.99 (1%), -3.43 (5%), -3.14 (10%)
Null Hypothesis: The process contains a unit root.
Alternative Hypothesis: The process is weakly stationary.

Cointegration Analysis (Johansen)

```
In [4]: # Cointegration Analysis (Johansen)

johansen_test = coint_johansen(data[['ln_pce', 'ln_dpi']], det_order=1, k_ar_diff=2)
print("== Johansen Test ==")
print("Trace statistic:", johansen_test.trace_stat)
print("Critical values:", johansen_test.cvt)

== Johansen Test ==
Trace statistic: [19.4704735  3.65278798]
Critical values: [[16.1619 18.3985 23.1485]
 [ 2.7055  3.8415  6.6349]]
```

VECM Estimation

```
In [5]: # 4. VECM Estimation

vecm_model = VECM(
    data[['ln_pce', 'ln_dpi']],
    k_ar_diff=4,
    coint_rank=1,
    deterministic="ci"
)
vecm_fit = vecm_model.fit()
print("== VECM Summary ==")
print(vecm_fit.summary())
```

==== VECM Summary ===

Det. terms outside the coint. relation & lagged endog. parameters for equation ln_pce

	coef	std err	z	P> z	[0.025	0.975]
<hr/>						
L1.ln_pce	-0.0261	0.065	-0.402	0.688	-0.153	0.101
L1.ln_dpi	0.3358	0.050	6.774	0.000	0.239	0.433
L2.ln_pce	0.0471	0.064	0.741	0.459	-0.077	0.172
L2.ln_dpi	0.1633	0.050	3.273	0.001	0.065	0.261
L3.ln_pce	0.0067	0.065	0.104	0.917	-0.120	0.134
L3.ln_dpi	0.1225	0.049	2.517	0.012	0.027	0.218
L4.ln_pce	0.0653	0.067	0.977	0.329	-0.066	0.196
L4.ln_dpi	0.0149	0.048	0.307	0.759	-0.080	0.110

Det. terms outside the coint. relation & lagged endog. parameters for equation ln_dpi

	coef	std err	z	P> z	[0.025	0.975]
<hr/>						
L1.ln_pce	-0.0009	0.077	-0.011	0.991	-0.153	0.151
L1.ln_dpi	-0.1280	0.059	-2.167	0.030	-0.244	-0.012
L2.ln_pce	0.3102	0.076	4.100	0.000	0.162	0.458
L2.ln_dpi	0.0338	0.059	0.569	0.570	-0.083	0.150
L3.ln_pce	-0.4265	0.077	-5.528	0.000	-0.578	-0.275
L3.ln_dpi	0.1596	0.058	2.752	0.006	0.046	0.273
L4.ln_pce	-0.0341	0.080	-0.428	0.668	-0.190	0.122
L4.ln_dpi	-0.0357	0.058	-0.618	0.537	-0.149	0.078

Loading coefficients (alpha) for equation ln_pce

	coef	std err	z	P> z	[0.025	0.975]
<hr/>						
ec1	0.0137	0.009	1.572	0.116	-0.003	0.031

Loading coefficients (alpha) for equation ln_dpi

	coef	std err	z	P> z	[0.025	0.975]
<hr/>						
ec1	0.0672	0.010	6.473	0.000	0.047	0.088

Cointegration relations for loading-coefficients-column 1

	coef	std err	z	P> z	[0.025	0.975]
<hr/>						
beta.1	1.0000	0	0	0.000	1.000	1.000
beta.2	-1.0827	0.018	-60.776	0.000	-1.118	-1.048





```
const      0.9904      0.166      5.982      0.000      0.666      1.315
=====
```


Post-Estimation Analysis (IRF)

```
In [6]: # Post-Estimation Analysis (IRF)
```

```
irf = vecm_fit.irf(20)
irf.plot(orth=True, signif=0.05)
plt.suptitle("Impulse Response Function (PCE & DPI)")
plt.show()
```

Impulse Response Function (PCE & DPI)





Forecasting with Confidence Intervals

```
In [7]: # 6. Forecasting with Confidence Intervals

steps = 12 # forecast horizon
forecast_vals, lower_vals, upper_vals = vecm_fit.predict(steps=steps, alpha=0.05)

forecast_df = pd.DataFrame(np.exp(forecast_vals), columns=['PCE_forecast', 'DPI_forecast'])
lower_df = pd.DataFrame(np.exp(lower_vals), columns=['PCE_lower', 'DPI_lower'])
upper_df = pd.DataFrame(np.exp(upper_vals), columns=['PCE_upper', 'DPI_upper'])

forecast_ci = pd.concat([forecast_df, lower_df, upper_df], axis=1)
print("VECM Forecast with 95% CI: \n")
forecast_ci
```

VECM Forecast with 95% CI:

Out[7]:

	PCE_forecast	DPI_forecast	PCE_lower	DPI_lower	PCE_upper	DPI_upper
0	16543.509330	18112.374014	16238.812727	17715.737180	16853.923100	18517.891131
1	16619.298163	18252.004211	16188.903973	17743.779790	17061.134706	18774.785399
2	16708.520962	18344.364848	16162.878490	17734.628591	17272.583772	18975.064515
3	16791.501662	18418.482711	16122.917170	17680.559691	17487.810989	19187.203985
4	16873.389677	18531.452256	16096.125862	17727.149409	17688.186687	19372.247325
5	16958.663320	18617.822729	16080.875221	17749.370501	17884.366220	19528.767128
6	17039.671355	18710.325189	16067.704249	17785.179251	18070.434668	19683.595185
7	17122.509174	18809.794830	16062.234787	17840.367128	18252.772686	19831.900264
8	17205.502602	18899.216071	16057.813410	17884.586199	18435.219804	19971.408011
9	17287.070353	18994.797173	16053.403249	17938.453279	18615.541936	20113.346118
10	17369.573818	19089.078886	16049.905333	17991.212318	18797.749167	20253.939883
11	17451.527181	19180.963538	16044.696639	18040.810430	18981.711391	20393.172673


Plot forecast

In [8]:

```
# Plot forecast
last_date = data.index[-1]
future_index = pd.date_range(start=last_date, periods=steps+1, freq="M")[1:]

plt.figure(figsize=(10,6))
plt.plot(data.index, data['PCE'], label='Historical PCE')
plt.plot(future_index, forecast_df['PCE_forecast'], color='red', label='PCE Forecast')
plt.fill_between(future_index, lower_df['PCE_lower'], upper_df['PCE_upper'], color='red', alpha=0.2, label='95% CI')
plt.title('PCE Forecast with 95% CI')
plt.legend()
plt.show()
```

PCE Forecast with 95% CI



Compare VECM to VAR

```
In [9]: # Compare VECM to VAR
var_model = VAR(data[['ln_pce', 'ln_dpi']])
var_fit = var_model.fit(maxlags=4)
var_forecast = var_fit.forecast(y=data[['ln_pce', 'ln_dpi']].values[-var_fit.k_ar:], steps=12)
var_forecast_df = pd.DataFrame(np.exp(var_forecast), columns=['PCE_forecast_VAR', 'DPI_forecast_VAR'])
```

```
print("VAR Forecast (Next 12 Periods): \n")
var_forecast_df
```

VAR Forecast (Next 12 Periods):

Out[9]:

	PCE_forecast_VAR	DPI_forecast_VAR
0	16548.662420	18108.083644
1	16630.219789	18246.313860
2	16732.309389	18334.339971
3	16824.305692	18405.481622
4	16914.942806	18517.257608
5	17010.024206	18599.808467
6	17100.371113	18690.859354
7	17192.581364	18790.163547
8	17285.406027	18879.888526
9	17377.210400	18976.696839
10	17470.364657	19073.445060
11	17563.525934	19168.533887

Include GDP as additional variable in VECM

In [10]:

```
# Include GDP as additional variable in VECM
vecm_model_gdp = VECM(
    data[['ln_pce', 'ln_dpi', 'ln_gdp']],
    k_ar_diff=4,
    coint_rank=1,
    deterministic="ci"
)
vecm_fit_gdp = vecm_model_gdp.fit()
```

```
print("== VECM with GDP Summary ==")
print(vecm_fit_gdp.summary())
```

==== VECM with GDP Summary ===

Det. terms outside the coint. relation & lagged endog. parameters for equation ln_pce

	coef	std err	z	P> z	[0.025	0.975]
<hr/>						
L1.ln_pce	-0.0609	0.094	-0.649	0.516	-0.245	0.123
L1.ln_dpi	0.2871	0.053	5.392	0.000	0.183	0.391
L1.ln_gdp	0.0173	0.083	0.209	0.835	-0.145	0.180
L2.ln_pce	-0.0177	0.091	-0.193	0.847	-0.197	0.161
L2.ln_dpi	0.1140	0.052	2.202	0.028	0.013	0.215
L2.ln_gdp	0.0445	0.081	0.550	0.582	-0.114	0.203
L3.ln_pce	0.1235	0.089	1.391	0.164	-0.050	0.297
L3.ln_dpi	0.0614	0.050	1.232	0.218	-0.036	0.159
L3.ln_gdp	-0.1757	0.081	-2.158	0.031	-0.335	-0.016
L4.ln_pce	0.1614	0.090	1.795	0.073	-0.015	0.338
L4.ln_dpi	-0.0263	0.048	-0.545	0.586	-0.121	0.068
L4.ln_gdp	-0.1388	0.081	-1.712	0.087	-0.298	0.020

Det. terms outside the coint. relation & lagged endog. parameters for equation ln_dpi

	coef	std err	z	P> z	[0.025	0.975]
<hr/>						
L1.ln_pce	0.1779	0.117	1.518	0.129	-0.052	0.408
L1.ln_dpi	-0.1596	0.066	-2.401	0.016	-0.290	-0.029
L1.ln_gdp	-0.1232	0.103	-1.191	0.234	-0.326	0.080
L2.ln_pce	0.1461	0.114	1.282	0.200	-0.077	0.369
L2.ln_dpi	0.0163	0.065	0.252	0.801	-0.110	0.143
L2.ln_gdp	0.2862	0.101	2.832	0.005	0.088	0.484
L3.ln_pce	-0.3477	0.111	-3.138	0.002	-0.565	-0.131
L3.ln_dpi	0.1263	0.062	2.030	0.042	0.004	0.248
L3.ln_gdp	-0.0770	0.102	-0.757	0.449	-0.276	0.122
L4.ln_pce	0.1646	0.112	1.466	0.143	-0.055	0.385
L4.ln_dpi	-0.0518	0.060	-0.860	0.390	-0.170	0.066
L4.ln_gdp	-0.1712	0.101	-1.691	0.091	-0.370	0.027

Det. terms outside the coint. relation & lagged endog. parameters for equation ln_gdp

	coef	std err	z	P> z	[0.025	0.975]
<hr/>						
L1.ln_pce	0.0047	0.108	0.044	0.965	-0.207	0.216
L1.ln_dpi	0.2705	0.061	4.418	0.000	0.151	0.391
L1.ln_gdp	0.1247	0.095	1.308	0.191	-0.062	0.312
L2.ln_pce	-0.1807	0.105	-1.721	0.085	-0.387	0.025
L2.ln_dpi	-0.0039	0.060	-0.065	0.948	-0.121	0.113

L2.ln_gdp	0.2541	0.093	2.727	0.006	0.071	0.437
L3.ln_pce	-0.0920	0.102	-0.902	0.367	-0.292	0.108
L3.ln_dpi	0.0216	0.057	0.377	0.706	-0.091	0.134
L3.ln_gdp	-0.0162	0.094	-0.173	0.863	-0.200	0.167
L4.ln_pce	0.0990	0.103	0.957	0.339	-0.104	0.302
L4.ln_dpi	-0.0018	0.055	-0.032	0.974	-0.110	0.107
L4.ln_gdp	0.0667	0.093	0.715	0.475	-0.116	0.250
Loading coefficients (alpha) for equation ln_pce						
=====						
	coef	std err	z	P> z	[0.025	0.975]

ec1	-0.0232	0.006	-4.030	0.000	-0.035	-0.012
Loading coefficients (alpha) for equation ln_dpi						
=====						
	coef	std err	z	P> z	[0.025	0.975]

ec1	-0.0280	0.007	-3.891	0.000	-0.042	-0.014
Loading coefficients (alpha) for equation ln_gdp						
=====						
	coef	std err	z	P> z	[0.025	0.975]


ec1	-0.0273	0.007	-4.122	0.000	-0.040	-0.014
Cointegration relations for loading-coefficients-column 1						
=====						
	coef	std err	z	P> z	[0.025	0.975]


beta.1	1.0000	0	0	0.000	1.000	1.000
beta.2	-0.9194	0.221	-4.166	0.000	-1.352	-0.487
beta.3	-0.0051	0.106	-0.048	0.962	-0.213	0.202
const	-0.8480	1.074	-0.790	0.430	-2.953	1.257
=====						

IRF with GDP

```
In [11]: # IRF with GDP
irf_gdp = vecm_fit_gdp.irf(20)
irf_gdp.plot(orth=True, signif=0.05)
plt.suptitle("Impulse Response Function (PCE, DPI, GDP)")
plt.show()
```

Impulse Response Function (PCE, DPI, GDP)





Forecast with GDP

In [12]:

```
# Forecast with GDP
forecast_gdp_vals, lower_gdp_vals, upper_gdp_vals = vecm_fit_gdp.predict(steps=12, alpha=0.05)
forecast_gdp_df = pd.DataFrame(np.exp(forecast_gdp_vals), columns=['PCE_forecast', 'DPI_forecast', 'GDP_forecast'])
print("VECM Forecast with GDP (Next 12 Periods): \n")
forecast_gdp_df
```

VECM Forecast with GDP (Next 12 Periods):

Out[12]:

	PCE_forecast	DPI_forecast	GDP_forecast
0	16564.106694	18117.732846	30851.581843
1	16657.127985	18289.149455	31211.855568
2	16751.435271	18374.981808	31566.031606
3	16835.808606	18443.685209	31910.858139
4	16923.823736	18556.024800	32253.629708
5	17010.903171	18637.144105	32611.197661
6	17096.139823	18730.557103	32956.103865
7	17184.857530	18829.080822	33309.847955
8	17272.784003	18919.387898	33665.370745
9	17360.869605	19015.968361	34021.558607
10	17450.569808	19111.185668	34381.515074
11	17540.054615	19205.932884	34744.023063

In [12]: